

History of Boosting Algorithm: Adaboost, GBM, XGBoost

Seungyeop Hyun

Seoul National University

July 28, 2022

Contents

1 Introduction

2 AdaBoost

3 Gradient Boosting

4 XGBoost

Boosting Algorithm

- Boosting algorithm is proposed by Freund and Schapire(1996).
- Boosting is an ensemble method that combines multiple weak learners into a strong learner.
- It trains each weak learner sequentially, while the bagging uses independent models.
- Weak learner is one whose error rate is slightly better than random guessing.

Boosting Algorithm

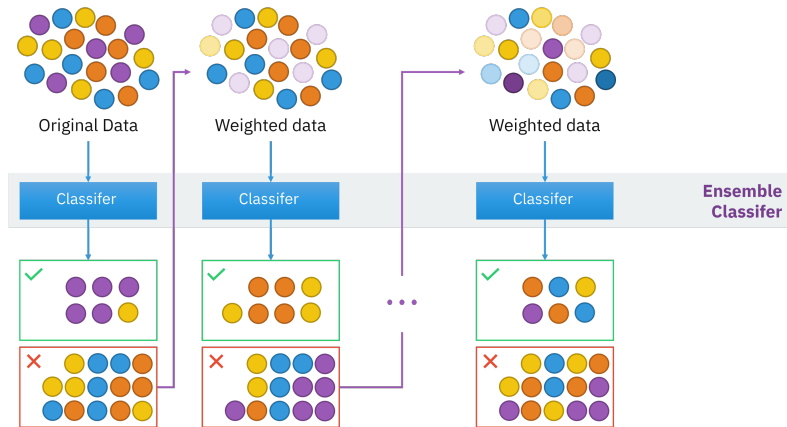


Figure 1: An illustration presenting the intuition behind the boosting algorithm, consisting of the parallel learners and weighted dataset.

Contents

1 Introduction

2 AdaBoost

3 Gradient Boosting

4 XGBoost

AdaBoost

- Freund and Schapire proposed the adaptive boosting, called AdaBoost in 1996.
- AdaBoost repeatedly updates weights and learns the weak learner using modified versions of the data in every iteration.
- AdaBoost uses the stump as weak learner.



Schematic of AdaBoost

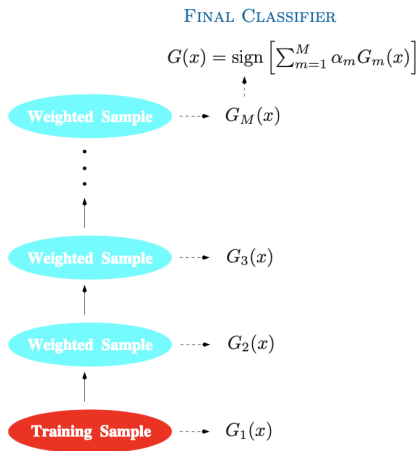


Figure 2: Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.

Algorithm of AdaBoost

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

Algorithm of AdaBoost

- Observations that were misclassified by the weak learner $G_{m-1}(x)$ at the previous step have increased weights, whereas the weights are decreased for those that were classified correctly.
- $G_m(x)$ is fitted using weighted data from the previous step.
- The weighted error rate is calculated with the classifier $G_m(x)$.
- α_m is used for updating weights. And it works for the weight of each weak learner when we make the final classifier.

Why it works?

- We can understand AdaBoost with statistical learning mechanism.
- From the classifier $G(x) = \text{sign}[\sum_m^M \alpha_m G_m(x)]$, we can think the boosting as a way of fitting an additive expansion in a set of basis function.
- Basis function expansions take the form

$$f(x) = \sum_{m=1}^M \beta_m b(x; \gamma_m)$$

where β_m are the coefficients for basis function $b(x; \gamma) \in \mathbb{R}$, characterized by γ_m

- Here the $b(x; \gamma)$ are the individual learner $G_m(x) \in \{-1, 1\}$

Forward Stagewise Additive Modeling

- We can find the solution with forward stagewise additive modeling algorithm.
- It finds the answer by sequentially adding new basis functions to the expansion with the functions that have already been added are fixed.

Algorithm 10.2 *Forward Stagewise Additive Modeling.*

1. Initialize $f_0(x) = 0$.

2. For $m = 1$ to M :

(a) Compute

$$(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma)).$$

(b) Set $f_m(x) = f_{m-1}(x) + \beta_m b(x; \gamma_m)$.

Equivalence between two algorithms

- We can see that AdaBoost is equivalent to forward stagewise additive modeling using the 'exponential' loss function

$$L(y, f(x)) = \exp(-yf(x))$$

- We have to solve

$$(\beta_m, G_m) = \arg \min_{\beta, G} \sum_{i=1}^n \exp(-y_i(f_{m-1}(x_i) + \beta G(x_i)))$$

and we can rewrite it with $w_i^{(m)} = \exp(-y_i f_{m-1}(x_i))$,

$$(\beta_m, G_m) = \arg \min_{\beta, G} \sum_{i=1}^n w_i^{(m)} \exp(-\beta y_i G(x_i)) \quad (2.1)$$

Equivalence between two algorithms

- For fixed $\beta > 0$, G_m in (2.1) is calculated by

$$G_m = \arg \min_G \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i))$$

and it minimizes the weighted error rate

$$\text{err}_m = \frac{\sum_{i=1}^N w_i^{(m)} I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i^{(m)}}$$

- We can get the solution of β by plugging G_m into (2.1)

$$\beta_m = \frac{1}{2} \log \frac{1 - \text{err}_m}{\text{err}_m}$$

Equivalence between two algorithms

- Then the $f_m(x)$ is updated

$$f_m(x) = f_{m-1}(x) + \beta_m G_m(x)$$

- And the weights for the next step automatically changed

$$w_i^{(m+1)} = w_i^{(m)} e^{-\beta_m y_i G_m(x_i)} \quad (2.2)$$

- With the fact that $-y_i G_m(x_i) = 2 \cdot I(y_i \neq G_m(x_i)) - 1$, (2.2) becomes

$$w_i^{(m+1)} = w_i^{(m)} e^{\alpha_m I(y_i \neq G_m(x_i))} \cdot e^{-\beta_m} \quad (2.3)$$

where $\alpha_m = 2\beta_m$ from AdaBoost.

Equivalence between two algorithms

- Because the $e^{-\beta_m}$ is multiplied to all the weights by same value, it has no effect on fitting the classifier.
- The weight update step in AdaBoost is equivalent to (2.3).
- We can see that the AdaBoost gets an answer by solving the optimization problem on loss functions.

Contents

- 1 Introduction
- 2 AdaBoost
- 3 Gradient Boosting**
- 4 XGBoost

Introduction

- Forward stagewise boosting(or AdaBoost) is a very greedy strategy that follows the optimal direction to minimize loss function.
- However, we can't find a solution when we don't use some loss functions like exponential loss.
- Friedman (2001) extended AdaBoost for general loss functions which are differentiable.
- Purpose doesn't change. We want to solve the problem below.

$$(\beta_m, \gamma_m) = \arg \min_{\gamma, \beta} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta h(x_i; \gamma))$$

Gradient Boosting

- Gradient boosting uses the gradient of loss function

$$g_m(x_i) = \left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right]_{F=F_{m-1}}$$

- But the best steepest descent direction $-\mathbf{g}_m = \{-g_m(x_i)\}_1^N$ is defined only at the data points $\{x_i\}_1^N$.
- Gradient boosting chooses the basis function $h(x_i; \gamma_m)$ that produces $\mathbf{h}_m = \{h(x_i; \gamma_m)\}$ most parallel to $-\mathbf{g}_m \in \mathbb{R}^N$.

$$\gamma_m = \arg \min_{\gamma, \beta} \sum_{i=1}^n [-g_m(x_i) - \beta h(x_i; \gamma)]^2$$

Algorithm of Gradient Boosting

Algorithm 1: Gradient_Boost

```
1   $F_0(\mathbf{x}) = \arg \min_{\rho} \sum_{i=1}^N L(y_i, \rho)$ 
2  For  $m = 1$  to  $M$  do:
3       $\tilde{y}_i = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(\mathbf{x})=F_{m-1}(\mathbf{x})}, i = 1, N$ 
4       $\mathbf{a}_m = \arg \min_{\mathbf{a}, \beta} \sum_{i=1}^N [\tilde{y}_i - \beta h(\mathbf{x}_i; \mathbf{a})]^2$ 
5       $\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(\mathbf{x}_i) + \rho h(\mathbf{x}_i; \mathbf{a}_m))$ 
6       $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \rho_m h(\mathbf{x}; \mathbf{a}_m)$ 
7  endFor
end Algorithm
```

Application: Gradient Tree Boosting

- Basis function h , regression tree with J terminal nodes, can be expressed the additive form

$$h(x; \{b_j, R_j\}_1^J) = \sum_{j=1}^J w_j I(x \in R_j)$$

where w_j is the weight for the terminal node R_j .

- From the gradient boosting algorithm, the model update step is

$$\begin{aligned} f_m(x) &= f_{m-1}(x) + \rho_m \sum_{j=1}^J w_{jm} I(x \in R_{jm}) \\ &= f_{m-1}(x) + \sum_{j=1}^J \gamma_{jm} I(x \in R_{jm}) \end{aligned}$$

where $\gamma_{jm} = \rho_m w_{jm}$

Application: Gradient Tree Boosting

- We can optimize γ_{jm}

$$(\gamma_{1m}, \dots, \gamma_{Jm}) = \arg \min_{\gamma_{1m}, \dots, \gamma_{Jm}} \sum_{i=1}^N L(y_i, F_{m-1}(x_i) + \sum_{j=1}^J \gamma_j I(x \in R_{jm}))$$

$$\iff \gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma), \quad j = 1, \dots, J$$

Contents

1 Introduction

2 AdaBoost

3 Gradient Boosting

4 XGBoost

eXtreme Gradient Boosting

- Tianqi Chen and Carlos Guestrin(2016) proposed XGBoost, eXtreme Gradient Boosting, which is based on gradient boosting.
- XGboost is faster and more accurate than gradient boosting algorithm.
- It improves overfitting problem of gradient boosting.

Tree Ensemble Models

- For a data set with n examples and m features

$\mathcal{D} = \{(\mathbf{x}_i, y_i)\}$ ($|\mathcal{D}| = n$, $\mathbf{x}_i \in \mathbb{R}^m$, $y_i \in \mathbb{R}$), a tree ensemble model uses K additive functions to predict the output.

$$\hat{y}_i = \phi(\mathbf{x}_i) = \sum_{k=1}^K f_k(\mathbf{x}_i), \quad f_k \in \mathcal{F}$$

where $\mathcal{F} = \{f(\mathbf{x}) = w_{q(\mathbf{x})}\} (q : \mathbb{R}^m \rightarrow T, w \in \mathbb{R}^T)$ is the space of basic tree model.

Tree Ensemble Models

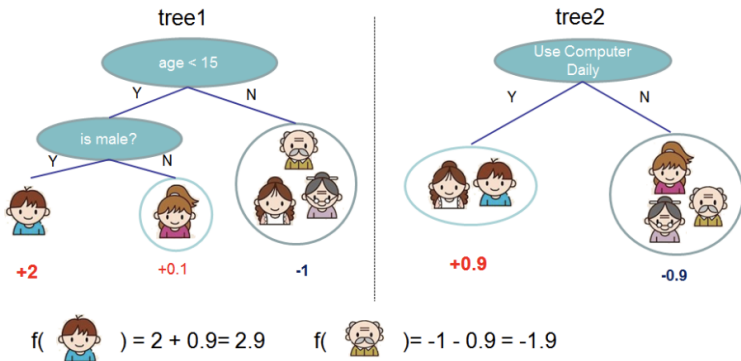


Figure 3: Tree Ensemble Model. The final prediction for a given example is the sum of predictions from each tree.

Regularized Learning Objective

- To learn the individual trees in XGBoost, we minimize the *regularized* objective.

$$\mathcal{L}^{(t)}(\phi) = \sum_i l(y_i, \hat{y}_i) + \Omega(f_t)$$

$$\text{where } \Omega(f_t) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$

- l is a differentiable convex loss function.
- Ω penalized the complexity of the trees. It helps to make simple functions and smooth the weights for leaf nodes to avoid overfitting.
- If the second term Ω goes to 0, it is equal to objective used in traditional gradient boosting.

Gradient Tree Boosting

- In XGBoost, the model is trained in an additive manner. We learn the each tree sequentially.
- Let the $\hat{y}_i^{(t)}$ be the prediction value of the i -th examples at the t -th iteration. We have to find optimal f_t in below,

$$\mathcal{L}^{(t)}(\phi) = \sum_i l(y_i, \hat{y}_i^{(t-1)} + f_t(\mathbf{x}_i)) + \Omega(f_t)$$

- We greedily add the f_t which minimizes the objective.

Gradient Tree Boosting

- Using second order approximation, we can quickly optimize the objective.

$$\mathcal{L}^{(t)} \simeq \sum_{i=1}^n \left[l(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(\mathbf{x}_i) + \frac{1}{2} h_i f_t^2(\mathbf{x}_i) \right] + \Omega(f_t)$$

where $g_i = \partial_{\hat{y}^{(t-1)}} l(y_i, \hat{y}_i^{(t-1)})$, $h_i = \partial_{\hat{y}^{(t-1)}}^2 l(y_i, \hat{y}_i^{(t-1)})$ are the first and second derivative of loss function.

- And we can erase the constant terms to simplify the objective at step t .

$$\tilde{\mathcal{L}}^{(t)} = \sum_{i=1}^n \left[g_i f_t(\mathbf{x}_i) + \frac{1}{2} h_i f_t^2(\mathbf{x}_i) \right] + \Omega(f_t)$$

Gradient Tree Boosting

- Set $I_j = \{i | q(\mathbf{x}_i) = j\}$ as the observation set of the leaf j . The equation can be rewritten By expanding Ω

$$\begin{aligned}\tilde{\mathcal{L}}^{(t)} &= \sum_{i=1}^n \left[g_i f_t(\mathbf{x}_i) + \frac{1}{2} h_i f_t^2(\mathbf{x}_i) \right] + \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2 \\ &= \sum_{j=1}^T \left[\left(\sum_{i \in I_j} g_i \right) w_j + \frac{1}{2} \left(\sum_{i \in I_j} h_i + \lambda \right) w_j^2 \right] + \gamma T\end{aligned}$$

Gradient Tree Boosting

- Then we can compute the optimal weight and value of the simplified objective.

$$w_j^* = -\frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda}$$
$$\tilde{\mathcal{L}}^{(t)}(q) = -\frac{1}{2} \sum_{j=1}^T \frac{(\sum_{i \in I_j} g_i)^2}{\sum_{i \in I_j} h_i + \lambda} + \gamma T$$

- $\tilde{\mathcal{L}}^{(t)}(q)$ used as a scoring function to measure the quality of a tree structure q .
- At an any single node, the loss reduction after the split is given By

$$\mathcal{L}_{split} = \frac{1}{2} \left[\frac{(\sum_{i \in I_L} g_i)^2}{\sum_{i \in I_L} h_i + \lambda} + \frac{(\sum_{i \in I_R} g_i)^2}{\sum_{i \in I_R} h_i + \lambda} - \frac{(\sum_{i \in I} g_i)^2}{\sum_{i \in I} h_i + \lambda} \right] - \gamma$$

where I_L, I_R are the observation sets of separated nodes.

Shrinkage and Column Subsampling

- Shrinkage and column subsampling are two additional techniques to prevent overfitting.
- Shrinkage factor η is multiplied to each tree model. It reduces the influence of each tree and gives a chance to improve the whole model for the trees at future steps.
- Column subsampling presents randomness for learning each tree. It is from Random Forest.

Split Finding Algorithm

- Finding the best split is important issue in tree learning.
- Exact greedy algorithm enumerates over all the possible splits on all the features.

Algorithm 1: Exact Greedy Algorithm for Split Finding

Input: I , instance set of current node

Input: d , feature dimension

$gain \leftarrow 0$

$G \leftarrow \sum_{i \in I} g_i, H \leftarrow \sum_{i \in I} h_i$

for $k = 1$ **to** m **do**

$G_L \leftarrow 0, H_L \leftarrow 0$

for j in sorted(I , by \mathbf{x}_{jk}) **do**

$G_L \leftarrow G_L + g_j, H_L \leftarrow H_L + h_j$

$G_R \leftarrow G - G_L, H_R \leftarrow H - H_L$

$score \leftarrow \max(score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda})$

end

end

Output: Split with max score

Approximate Algorithm

- Exact greedy algorithm is powerful, but it is computationally impossible to find all candidates on continuous features.

Algorithm 2: Approximate Algorithm for Split Finding

for $k = 1$ **to** m **do**

 Propose $S_k = \{s_{k1}, s_{k2}, \dots, s_{kl}\}$ by percentiles on feature k .
 Proposal can be done per tree (global), or per split(local).

end

for $k = 1$ **to** m **do**

$G_{kv} \leftarrow \sum_{j \in \{j | s_{k,v} \geq \mathbf{x}_{jk} > s_{k,v-1}\}} g_j$
 $H_{kv} \leftarrow \sum_{j \in \{j | s_{k,v} \geq \mathbf{x}_{jk} > s_{k,v-1}\}} h_j$

end

Follow same step as in previous section to find max score only among proposed splits.

Approximate Algorithm

- The algorithm first suggests candidate split points according to percentiles of feature distribution through weighted quantile sketch algorithm.
- Then, it puts the features into the groups split by the candidate points calculated in advance.
- There are two variants, global and local variant.
 - ① Global : proposing all the candidate splits at the initial phase of learning.
 - ② Local : re-proposing the points after each split.

Approximate Algorithm

- Global variant needs more candidates because the split points are not changed during the training. And the local variant is more proper for deeper trees.




Figure 4: Global and local proposal.

Approximate Algorithm




Figure 5: The eps parameter corresponds to the accuracy of the approximate sketch. This roughly translates to $1 / \text{eps}$ buckets in the proposal. We find that local proposals require fewer buckets, because it refines split candidates

Weighted Quantile Sketch

- Proposing the candidate split points is an important step in the approximate algorithm. XGBoost uses *weighted quantile sketch* algorithm in this stage.
- Let $\mathcal{D}_k = \{(x_{1k}, h_1), \dots, (x_{nk}, h_k)\}$ means k -th feature values and second order gradient statistics of each training examples.
- Define the rank functions $r_k : \mathbb{R} \rightarrow [0, +\infty)$ as

$$r_k(z) = \frac{1}{\sum_{(x,h) \in \mathcal{D}_k} h} \sum_{(x,h) \in \mathcal{D}_k, x < z} h$$

- It represents the proportion of instances.

Weighted Quantile Sketch

- Our goal is find candidate split points $\{s_{k1}, s_{k2}, \dots, s_{kl}\}$ satisfying

$$|r_k(s_{k,j}) - r_k(s_{k,j+1})| < \epsilon, \quad s_{k1} = \min_i \mathbf{x}_{ik}, s_{kl} = \max_i \mathbf{x}_{ik}$$

where ϵ is an approximation factor.

- Roughly $1/\epsilon$ candidate points are picked up.
- h_i is considered as weight of each data point. The objective at step t , $\tilde{\mathcal{L}}^{(t)}$, can be rewritten. It shows the role of h_i

$$\sum_{i=1}^n \frac{1}{2} h_i (f_t(\mathbf{x}_i) - g_i/h_i)^2 + \Omega(f_t) + \text{constant}$$

Sparsity-aware Split Finding

- Sparsity-aware split finding algorithm is used in XGBoost to handel a sparsity pattern in data.
- There are various causes for sparsity in real world problems.
 - ① presence of missing values
 - ② frequent zero entries
 - ③ artifacts of feature engineering such as one-hot encoding
- Using a default direction in each node, missing values are classified.

Sparsity-aware Split Finding

Algorithm 3: Sparsity-aware Split Finding

Input: I , instance set of current node

Input: $I_k = \{i \in I | x_{ik} \neq \text{missing}\}$

Input: d , feature dimension

Also applies to the approximate setting, only collect statistics of non-missing entries into buckets

$gain \leftarrow 0$

$G \leftarrow \sum_{i \in I} g_i, H \leftarrow \sum_{i \in I} h_i$

for $k = 1$ **to** m **do**

// enumerate missing value goto right

$G_L \leftarrow 0, H_L \leftarrow 0$

for j in sorted(I_k , ascent order by \mathbf{x}_{jk}) **do**

$G_L \leftarrow G_L + g_j, H_L \leftarrow H_L + h_j$

$G_R \leftarrow G - G_L, H_R \leftarrow H - H_L$

$score \leftarrow \max(score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda})$

end

// enumerate missing value goto left

$G_R \leftarrow 0, H_R \leftarrow 0$

for j in sorted(I_k , descent order by \mathbf{x}_{jk}) **do**

$G_R \leftarrow G_R + g_j, H_R \leftarrow H_R + h_j$

$G_L \leftarrow G - G_R, H_L \leftarrow H - H_R$

$score \leftarrow \max(score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda})$

end

end

Output: Split and default directions with max gain

Sparsity-aware Split Finding

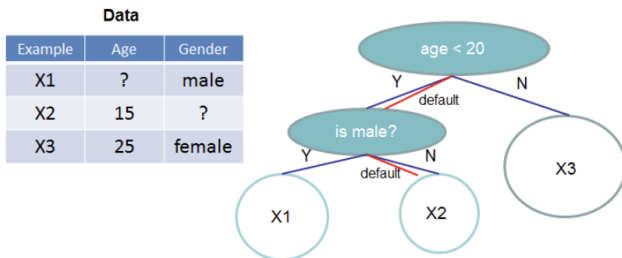


Figure 6: Tree structure with default directions. An example will be classified into the default direction when the feature needed for the split is missing.

Sparsity-aware Split Finding




Figure 7: Impact of the sparsity aware algorithm. The dataset is sparse mainly due to one-hot encoding.

- Recently, faster and stronger methods like LightGBM and CatBoost are proposed.
- LightGBM improves the learning speed dramatically and CatBoost works very well with categorical features.
- Boosting is still powerful comparing to the deep learning algorithms.

References

- Chen T and Guestrin C. XGBoost: A Scalable Tree Boosting System. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, p.785-794, 2016.
- Pilsung Kang. Business Analytics - IME654, Korea University
- Yoav Freund and Robert E. Schapire. a decision theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, p.119-139, 1996.
- Hastie T., Tibshirani R., Friedman J. H. (2001). The elements of statistical learning: Data mining, inference, and prediction.
- Friedman J. H.(2001), Greedy Function Approximation: A Gradient Boosting Machine, *The Annals of Statistics*